

Kevin Galim

Senior AI Research Engineer | Efficient LLM Inference, Generative Models, AI Systems
galimkevin@gmail.com · kevingalim.com · scholar.google.com · linkedin.com/in/kgalim/

Summary

Machine learning researcher specializing in efficient inference for large-scale generative models and LLM systems.

Author of 10+ publications at major ML conferences including ICLR, ICML, ACL, CVPR, and ECCV. Research focuses on scalable inference algorithms, speculative decoding, and efficient training and deployment of large generative models.

Experience

Senior AI Research Engineer

FuriosaAI

Jun 2021 – Present

Seoul, South Korea

- Conducted research on large-scale generative models and LLMs, including efficient inference (KV-cache optimization, diffusion LLMs), scalable training (PEFT, state space models), and advanced architectures.
- Co-authored multiple first- and co-first-author papers in top-tier conferences: ICLR, ICML, ACL, CVPR, ECCV, and WACV.
- Designed end-to-end pipelines for training and evaluating LLMs, including accelerator-aware rollout generation and inference pipelines on custom AI hardware.
- Built and demonstrated real-time computer vision demos on custom hardware at CVPR 2022 and 2023.

AI/Computer vision research and development

Funzin

Mar 2020 – Jun 2021

Seoul, South Korea

- Developed and trained deep learning models for object detection, road/sidewalk segmentation, and gesture detection using PyTorch and TensorFlow.
- Built perception systems for an autonomous golf cart platform, including detection and segmentation pipelines.
- Optimized models for embedded deployment using TensorRT, OpenVINO, Coral, DSP acceleration, and ARM NEON.
- Developed and calibrated a real-time OpenGL 3D surround-view system (SVM) for embedded hardware, demonstrated at CES 2021.

Technologies: PyTorch, TensorFlow, C++, Python, TensorRT, OpenVINO, ARM NEON

Web / AR Developer

Freelance

Aug 2019 – Mar 2020

Munich

- Built AR mobile applications using Unity3D.
- Developed cloud-backed web applications using AWS and Google Cloud.
- Created cross-platform mobile apps using Flutter.

Technologies: Unity3D, C#, JavaScript, Flutter, AWS, Google Cloud

C++/CUDA software engineer

ARRI

Aug 2015 – Sep 2016

Munich

- Developed GPU-accelerated image processing algorithms using CUDA and OpenCL.
- Built real-time visualization and image analysis tools using C++ and OpenGL.
- Contributed to software used in professional digital cinema workflows.

Technologies: C++, CUDA, OpenGL, OpenCL

Education

Technical University of Munich
Master's degree, Informatics: Games Engineering

2016 – 2019
Grade: 1.4 (German system)

The University of Tokyo
Research, Computer Graphics

2017 – 2018

Technical University of Munich
Bachelor's degree, Informatics: Games Engineering

2013 – 2016
Grade: 1.6 (German system)

Publications

First-author / co-first-author publications

- **Kevin Galim***, Ethan Ewer*, Wonjun Kang, Minjae Lee, Hyung Il Koo, Kangwook Lee. Draft-based Approximate Inference for LLMs. *ICLR*, 2026.
- Wonjun Kang*, **Kevin Galim***, Seunghyuk Oh*, Minjae Lee, Yuchen Zeng, Shuibai Zhang, Coleman Hooper, Yuezhou Hu, Hyung Il Koo, Nam Ik Cho, et al. ParallelBench: Understanding the Trade-offs of Parallel Decoding in Diffusion LLMs. *ICLR*, 2026.
- Wonjun Kang*, **Kevin Galim***, Yuchen Zeng*, Minjae Lee, Hyung Il Koo, Nam Ik Cho. State-offset Tuning: State-based Parameter-Efficient Fine-Tuning for State Space Models. *ACL*, 2025.
- **Kevin Galim***, Wonjun Kang*, Yuchen Zeng*, Hyung Il Koo, Kangwook Lee. Parameter-Efficient Fine-Tuning of State Space Models. *ICML*, 2025.
- Wonjun Kang*, **Kevin Galim***, Hyung Il Koo, Nam Ik Cho. Counting Guidance for High Fidelity Text-to-Image Synthesis. *WACV (Oral)*, 2025.
- Wonjun Kang*, **Kevin Galim***, Hyung Il Koo. Eta Inversion: Designing an Optimal Eta Function for Diffusion-based Real Image Editing. *ECCV*, 2024.

Other publications

- Seunghyuk Oh, Minjae Lee, **Kevin Galim**, Minseo Kim, Hyung Koo, Wonjun Kang, Hanbaek Lyu, Kangwook Lee. Inference-Aligned SFT for Diffusion LLMs via Group-based Trajectory Sampling. *ICLR DeLTa Workshop*, 2026.
- Minjae Lee*, Wonjun Kang*, Byeongkeun Ahn, Christian Classen, **Kevin Galim**, Seunghyuk Oh, Minghao Yan, Hyung Il Koo, Kangwook Lee. TABED: Test-Time Adaptive Ensemble Drafting for Robust Speculative Decoding in LLMs. *EACL*, 2026.
- Wonjun Kang, Byeongkeun Ahn, Minjae Lee, **Kevin Galim**, Seunghyuk Oh, Hyung Il Koo, Nam Ik Cho. UNCAGE: Contrastive Attention Guidance for Masked Generative Transformers in Text-to-Image Generation. *IEEE Access*, 2026.
- Maxim Maximov, **Kevin Galim**, Laura Leal-Taixe. Focus on Defocus: Bridging the Synthetic to Real Domain Gap for Depth Estimation. *CVPR*, 2020.

Skills

Machine Learning: LLMs · Diffusion Models · Parameter-Efficient Fine-Tuning · State Space Models · Reinforcement Learning

Systems & Optimization: LLM Inference Optimization · CUDA · C++ · PyTorch

Computer Vision: Depth Estimation · Segmentation · Object Detection

Languages

English (Full professional proficiency) · German (Native or bilingual proficiency) · Korean (Professional working proficiency)